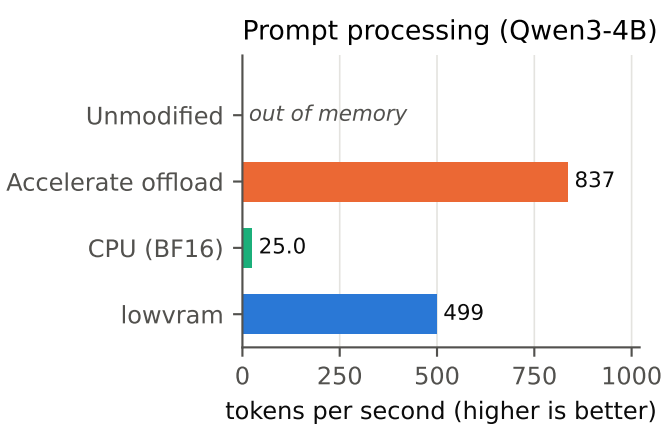
Prompt processing (Qwen3-4B)
Unmodified
out of memory
Accelerate offload
837
CPU (BF16)
25.0
lowvram
499
0
250
500
750
1000
tokens per second (higher is better)

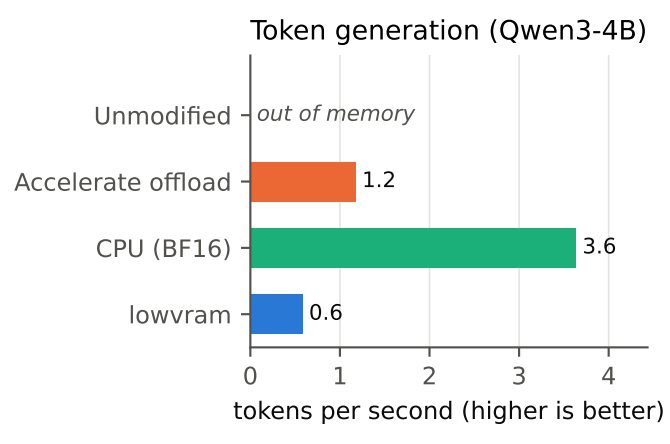
Token generation (Qwen3-4B)
Unmodified
out of memory
Accelerate offload
1.2
CPU (BF16)
3.6
lowvram
0.6
0
1
2
3
4
tokens per second (higher is better)